

Where to get AI news

The field moves weekly, and none of it reaches you through journals in time. Follow people, not hashtags.

- **X / Twitter** — still where things break first: model launches, papers, jailbreaks, the honest post-mortems.

@karpathy Andrej Karpathy (Anthropic; ex-OpenAI, ex-Tesla) — how the models actually work

@ylecun Yann LeCun (AMI Labs; ex-Meta) — the loudest sceptic on what LLMs cannot do

@rasbt Sebastian Raschka (independent; ex-Lightning AI) — implementation-level explainers, from scratch

@simonw Simon Willison (co-creator of Django) — tries every tool the day it ships

@steipete Peter Steinberger (OpenClaw) — agent setups and harness tricks

@sama Sam Altman (OpenAI) — the labs announce here first; read the rest as positioning

@DarioAmodei Dario Amodei (Anthropic; ex-OpenAI) — long essays on where he thinks this goes

@pesvklobouku Jakub Čížek (VTM.cz, dříve Živě.cz) — česky a srozumitelně

@kasikp Pavel Kasík (Seznam Zprávy) — česky, a s ověřováním hype

- **Newsletters**

The Batch (Andrew Ng, DeepLearning.AI) — plain weekly roundup, no hype dialect

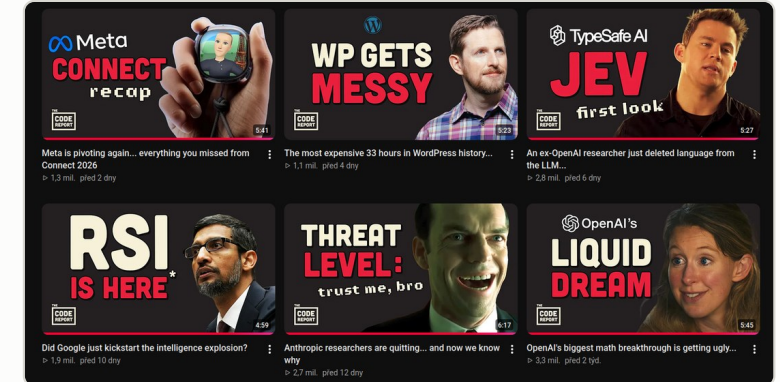
AI Snake Oil (Narayanan & Kapoor, Princeton) — academics taking the claims apart

- **Benchmarks**

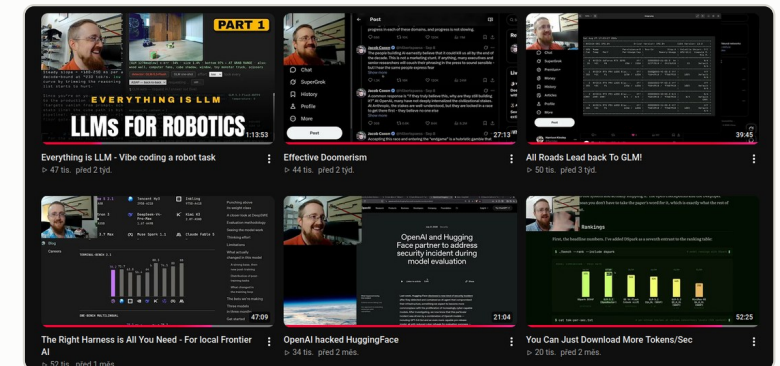
artificialanalysis.ai — quality, price and speed of every model side by side

tbench.ai — Terminal-Bench: agents doing real work in a terminal

Fireship — *The Code Report*: ~5 minutes, sarcastic, catches everything that happened this week.

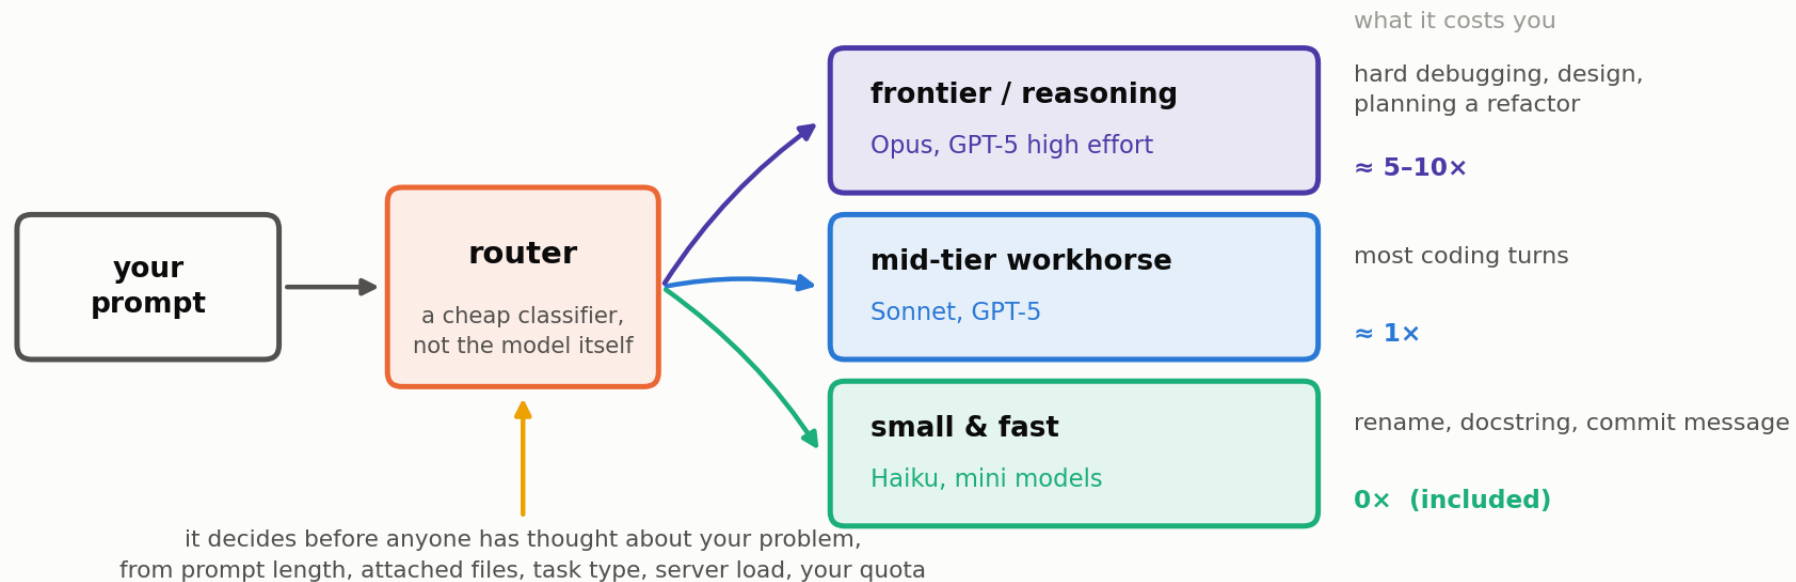


sentdex — long-form Python and ML: local models, harnesses, benchmarks run by hand.



Model routing — the “Auto” model

One prompt in — and something outside the model decides which model actually answers it.



- Copilot, Cursor and others offer **Auto**: you do not choose the model, a **router** does — per request. A cheap classifier **outside** the model, reading the surface of your prompt; it has not thought about your problem.
- Not the same router as in **MoE** (lecture 1): that one picks *experts inside one model*, per token, and was trained. This one picks *a whole model*, per request — a pricing decision.
- **Plan first — you will not one-shot a real problem.** Ask the strongest model for a numbered plan and nothing else (“do not change any files yet”), read it like a referee, keep it in `plan.md`, then spend one cheap prompt per step. That planning turn is where the 10× multiplier pays for itself.

Prompting tips

The model cannot see your screen, your data, or what you meant. Everything it knows about your problem, you put there.

vague

fix the plotting, it doesn't work

- ✗ which file? which plot?
- ✗ “doesn't work” — no error text
- ✗ so it guesses: four files edited, all of them confidently

specific

In hr_diagram.py the colour-magnitude plot is empty.

— the file, by name

Traceback: KeyError: 'phot_g_mean_mag', line 42.

— the real error, pasted

The csv is a Gaia DR3 gaia_source cone search.

— context it cannot guess

Fix the column lookup. Do not touch the query.

— what must not change

- **Write in English** — fewer tokens per word, better-trained behaviour — but keep local things local: Czech comments, labels and column names where they belong.
- **Drop the politeness.** “Could you please...”, “I was wondering if...” are tokens, not manners. Imperative, then the facts.
- **A question can be read as an order.** “Why is this loop so slow?” will get you a rewritten loop. If you only want an answer, say “**do not change any files, just explain**”.
- **Paste the real thing** — the command and the whole traceback, not your summary. Say what you already tried, or it will try it again.
- **Say what must not change** — “keep the query”, “do not touch the tests”. Everything you do not fence off is fair game.
- **One task per conversation.** Stale context from the previous problem is the most common cause of a stupid answer — start a new chat when you switch tasks.

Short projects

1 Open clusters with Gaia DR3

data: Gaia DR3 cone searches for **Pleiades**, **Praesepe** and **M67** (`astroquery.gaia`, table `gaia_source_lite`)

do: plot the proper motions, **cluster** the stars in (*pmra*, *pmdec*, *parallax*) to pick the members, then the HR diagram of each cluster and the age ordering from the turn-off

deliver: proper-motion diagram + HR diagram per cluster, member counts, distances, and the three clusters ordered by age

optional: isochrone fitting for a real age

2 Theoretical physics — to be announced

data: —

do: —

deliver: —